

## Prompt Pirates — HW2 Design Document

---

### 1. Error Mapping

Upstream (GitHub)	Our API	Rationale
201 (Created)	201 Created + Location: /issues/{number}	Correct REST semantics for creates.
200	200 OK	Happy path.
304 (If-None-Match)	304 Not Modified ( <i>optional extra credit</i> )	Conditional GET passthrough when ETag caching enabled.
401 / 403	401 Unauthorized	Token missing/invalid or repo access denied; normalize to 401 for clarity.
404	404 Not Found	Missing issue or resource.
422 (Validation)	400 Bad Request	Invalid input (state/title/body); map GitHub's validation error to client.
429 or secondary rate limit	429 Too Many Requests (or 503) + Retry-After	Surface backoff signal; caller can retry later.
5xx	502 Bad Gateway	Upstream fault; our gateway is healthy but GitHub failed.

---

### 2. Pagination Strategy

**Inputs:** state=open|closed|all, labels, page, per\_page (≤100).

**Behavior:**

- Forward page / per\_page directly to GitHub.
- Clamp per\_page to 100.
- Preserve GitHub's Link header so clients can follow rel="next" / rel="prev".

**Example:**

Link: <https://api.github.com/...&page=2>; rel="next",  
<...&page=4>; rel="last"

**Rationale:** Keeps client logic GitHub-compatible and avoids re-encoding pagination math.

---

### 3. Webhook Dedupe & Handling

- **Events:** ping, issues (opened/edited/closed/reopened), issue\_comment (created/edited/deleted).
- **Security:** HMAC SHA-256 over raw body with WEBHOOK\_SECRET; verify via hmac.compare\_digest; reject invalid requests with 401; never log raw secrets.
- **Idempotency:** Use X-GitHub-Delivery as a unique ID. Maintain in-memory seen\_ids; if duplicate → respond 204 immediately.
- **Processing:** Validate + record summary, respond 204 quickly; heavy work off-thread if necessary.

- **Debugging:** /events endpoint returns last N deliveries, e.g.:

```
[  
  { "id": "...", "event": "issue_comment", "action": "created",  
    "issue_number": 42, "timestamp": "..."}  
]
```

---

## 4. Security Trade-offs

### Chosen:

- Auth via fine-grained PAT (Issues: Read & Write, repo-scoped).
- Env vars (.env, ignored by git).
- Secrets masked in logs.
- Webhook HMAC with constant-time compare.
- CORS not broadly opened (service is server-to-server).

### Not Implemented (for balance):

- GitHub App (more secure, but more setup required).
- Persistent event store (in-memory buffer suffices).
- Role-based auth (not needed for single-consumer demo).

---

## 5. Reliability & Rate Limits

- **Backoff:** Inspect X-RateLimit-Remaining, X-RateLimit-Reset, Retry-After. On exhaustion → return 429/503 with helpful detail, forward Retry-After.
- **Timeouts:** httpx client defaults used. No auto-retry on writes (to preserve idempotency).
- **Health Check:** /healthz returns { "status": "ok" } (used by Docker/Compose).
- **Logs:** Structured, include path & upstream status; secrets always masked.

---

## 6. Data Models & OpenAPI

### Schemas:

- IssueIn: { title, body?, labels?[] }
- IssueOut: pass-through GitHub fields (number, html\_url, state, labels, timestamps)
- IssueUpdate: { title?, body?, state=open|closed }
- CommentIn: { body }
- CommentOut: { id, body, created\_at?, html\_url? }
- Error: { detail }

**Contract:** `openapi.yaml` (v3.1) defines all routes, params, examples, error models, and bearer security scheme.

---

## 7. Extra Credit

### ETag Support:

- Cache ETag/body from GET responses.
- On next request → send `If-None-Match`.
- If GitHub returns 304 → return 304 (or cached 200).
- Saves quota and bandwidth.

### Pipeline (GitHub Actions):

- **Stages:**
  1. **Lint** — `flake8` for Python style compliance
  2. **Test** — `pytest` with coverage ( $\geq 80\%$ )
  3. **Build** — Docker image build verification
  4. **Security** — Dependency vulnerability scan
- **Triggers:** push to `main`, pull requests.
- **Artifacts:** Coverage reports, Docker image (main branch only).
- **Environment:** Uses repo secrets for `GITHUB_TOKEN` (test scope).

---

## 8. Known Limitations / Future Work

- Webhook events lost on restart (could add SQLite/Redis).
- No background queue (all inline, though fast).
- PAT rotation is manual (GitHub App would improve).

---

## 9. Testing Summary

- **Unit Tests (pytest + respx):** happy & negative paths; webhook valid/invalid; pagination header; error mapping.
  - **Integration (with LIVE=1):** end-to-end flow — create → get → close/open → comment against real repo.
  - **Coverage:**  $\geq 80\%$  across app/ (achieved  $\sim 88\%$ ).
-